

Deep Learning Based Reverse Image Search System Using Feature Embeddings and Vector Similarity Search

A PROJECT REPORT

Submitted by

BHAVANA POLI - BL.SC.P2CSE25027

for the course
24CS739 – Deep Learning

Guided and Evaluated by

Dr. Santhanalakshmi
Dept. of CSE,



AMRITA SCHOOL OF ENGINEERING, BANGALORE
AMRITA VISHWA VIDHYAPEETHAM
BANGALORE-560 035
MAY 2025

Abstract

This report presents the design, implementation, and evaluation of a deep learning-based reverse image search system that retrieves visually similar images from a dataset using semantic feature embeddings and efficient vector similarity search. The system leverages a pretrained ResNet50 convolutional neural network to extract 2048-dimensional feature vectors from images, which are indexed using Facebook AI Similarity Search (FAISS) for scalable nearest-neighbour retrieval. The Caltech-101 benchmark dataset, comprising 101 object categories and approximately 9,000 images, is used for training and evaluation. The proposed approach achieves high retrieval accuracy, with Precision@10 of 0.87 and Mean Average Precision (MAP) of 0.83, while maintaining sub-second query response time. A Gradio-based interactive web interface enables real-time query submission and visual result display. The system demonstrates strong potential for deployment in e-commerce recommendation, medical imaging, and content-based digital asset management applications.

Keywords: *Content-Based Image Retrieval (CBIR), Deep Learning, ResNet50, FAISS, Feature Embeddings, Similarity Search, Gradio, Caltech-101*

1. Introduction

The exponential growth of digital image collections across social media platforms, e-commerce marketplaces, and medical imaging repositories has created an urgent demand for intelligent, content-aware image retrieval systems. Traditional keyword-based search engines rely on textual metadata that is often incomplete, inconsistent, or absent. Content-Based Image Retrieval (CBIR) systems address this limitation by enabling search directly over visual content, retrieving images that are semantically and visually similar to a given query image without requiring any textual description.

Classical CBIR systems relied on hand-crafted low-level features such as color histograms, edge maps, and texture descriptors (e.g., SIFT, SURF, HOG). While these methods are computationally efficient, they fail to capture high-level semantic relationships between images, leading to poor retrieval performance when query and target images differ in viewpoint, illumination, or scale.

The advent of deep convolutional neural networks (CNNs) has fundamentally transformed image representation learning. Pretrained models such as ResNet50, VGG16, and InceptionV3 generate rich, compact feature embeddings that encode semantic and structural information at multiple levels of abstraction. These deep features dramatically outperform hand-crafted descriptors on CBIR benchmarks.

However, performing exact nearest-neighbour search over large collections of high-dimensional feature vectors is computationally prohibitive. Facebook AI Similarity Search (FAISS) provides a solution through optimized approximate nearest-neighbour (ANN) indexing structures that scale to billions of vectors while maintaining sub-millisecond query times.

This project combines ResNet50 deep feature extraction with FAISS-based vector indexing to build a practical, end-to-end reverse image search system. The system is evaluated on the Caltech-101 benchmark dataset and exposes a Gradio-powered web interface for interactive querying.

1.1 Objectives

- Extract semantic feature embeddings from images using a pretrained ResNet50 model.
- Build a scalable FAISS vector index for fast approximate nearest-neighbour retrieval.
- Evaluate retrieval performance using standard CBIR metrics: Precision@K, Recall@K, MAP, MRR, and F1-Score.
- Deploy an interactive web interface for real-time visual search.
- Compare system performance against classical CBIR baselines.

2. Literature Review and Related Work

2.1 Classical CBIR Methods

Early content-based retrieval systems used low-level visual descriptors. Swain and Ballard (1991) proposed color histogram matching, while Lowe (2004) introduced SIFT keypoints for scale-invariant feature matching. Ojala et al. (2002) developed Local Binary Patterns (LBP) for texture-based retrieval. These methods are interpretable and fast but are sensitive to photometric and geometric variations.

2.2 Deep Learning for Image Retrieval

Babenko et al. (2014) demonstrated that CNN activations outperform traditional descriptors on landmark retrieval tasks. Razavian et al. (2014) showed that features from a pretrained AlexNet generalize remarkably well across diverse visual domains. He et al. (2016) introduced ResNet with residual connections, achieving state-of-the-art performance on ImageNet and producing highly discriminative embeddings for retrieval. Tolias et al. (2016) proposed Regional Maximum Activation of Convolutions (RMAC) for aggregating CNN features from multiple scales.

2.3 Approximate Nearest-Neighbour Search

Jegou et al. (2011) introduced Product Quantization (PQ) for compressed similarity search. Johnson et al. (2019) developed FAISS, which implements multiple ANN algorithms including Flat (exact), IVF (inverted file), PQ, and HNSW (hierarchical navigable small world). FAISS supports GPU acceleration and scales to billion-scale datasets.

2.4 End-to-End Retrieval Systems

Gordo et al. (2016) trained end-to-end systems with triplet loss for image retrieval. Noh et al. (2017) combined deep local features with geometric verification. Our system differs in that it uses off-the-shelf pretrained ResNet50 embeddings without fine-tuning, providing a baseline that is deployable without labeled triplet pairs.

3. Novelty and Contributions

The key novelty and contributions of this work are:

- Zero-shot visual search without fine-tuning: The system achieves strong CBIR performance using purely pretrained ResNet50 global average pooling features, demonstrating that off-the-shelf deep embeddings are sufficient for practical deployment.

- Integrated end-to-end pipeline: A seamless pipeline from raw image ingestion through preprocessing, feature extraction, FAISS indexing, similarity scoring, and interactive retrieval is implemented in a single deployable application.
- Lightweight deployment: The Gradio interface enables one-command deployment without cloud infrastructure, making the system accessible for small-scale production use.
- Reproducible benchmarking: The system is evaluated using five complementary retrieval metrics on the standard Caltech-101 benchmark, with results compared against three classical baselines.
- Persistent index storage: Feature embeddings, image paths, and FAISS index are persisted as .npy, .pkl, and .faiss files respectively, enabling incremental dataset updates without full re-indexing.

4. Dataset

4.1 Caltech-101

The primary dataset used in this project is Caltech-101, a widely used benchmark for image classification and retrieval. It was collected by Fei-Fei Li, Marco Andreetto, and Marc'Aurelio Ranzato at the California Institute of Technology.

Property	Details
Total images	~9,144
Number of categories	101 object classes + 1 background
Images per category	40 to 800 (mean ~58)
Image resolution	Approximately 300 x 200 pixels
Image format	JPEG
Category examples	Airplanes, Butterfly, Camera, Laptop, Elephant, Face, Motorbike
License	Free for non-commercial research use
Source	Caltech Vision Lab (http://www.vision.caltech.edu)

4.2 Dataset Split

For retrieval evaluation, the dataset is used as a gallery (database) of indexed images and a separate query set is constructed by selecting images not present in the gallery. A 80/20 gallery/query split is applied per category.

4.3 Challenges

- High intra-class variation due to different viewpoints, scales, and backgrounds.
- Low inter-class variance between visually similar categories (e.g., different aircraft types).
- Imbalanced class distribution with some categories having far fewer samples.

5. System Architecture

The system follows a two-phase pipeline: an offline indexing phase where all dataset images are processed and indexed, and an online query phase where a user-provided query image is processed and retrieved against the index in real time. Figure 1 illustrates the complete system architecture.




Figure 1: Complete System Architecture — Deep Learning Reverse Image Search

The architecture comprises eight major stages: (1) Query Image Input via Gradio UI, (2) Image Preprocessing, (3) Deep Feature Extraction via ResNet50, (4) 2048-dimensional Embedding Generation, (5) FAISS Vector Indexing with three persistent storage files, (6) Nearest-Neighbour Similarity Search, (7) Top-K Result Ranking, and (8) Result Display in the Gradio web interface.

6. Methodology

The methodology is divided into two phases: the Offline Indexing Phase executed once over the full dataset, and the Online Query Phase executed per user query at inference time.

6.1 Image Preprocessing

All input images undergo a standardized three-step preprocessing pipeline to satisfy the input requirements of the ResNet50 model and to ensure consistent feature extraction across all images.

Step 1 — Resize

Images are resized to a fixed spatial resolution of 224 x 224 pixels, matching the input dimensions expected by ResNet50.

```
I_resized = Resize(I_original, 224 x 224)
```

Step 2 — RGB Conversion

Grayscale images are converted to 3-channel RGB to satisfy the three-channel input requirement of ResNet50.

Step 3 — ImageNet Normalization

Pixel values are normalized using the ImageNet dataset statistics:

```
I_norm = (I_raw / 255 - mu) / sigma
mu      = [0.485, 0.456, 0.406]    (R, G, B channel means)
sigma   = [0.229, 0.224, 0.225]    (R, G, B channel std devs)
```

This normalization aligns the input distribution with that seen during ImageNet pretraining, preserving feature quality.

6.2 Deep Feature Extraction

Feature extraction is performed using a pretrained ResNet50 model from the Keras applications module. The model is configured without the top classification head and with global average pooling applied to the final convolutional feature map.

```
model = ResNet50(weights='imagenet', include_top=False, pooling='avg')
f(I) = GlobalAveragePool(Conv_50(I)) --> R^2048
```

The resulting feature vector $f(I)$ is a 2048-dimensional real-valued embedding that encodes the semantic content of the image: object structure, texture patterns, color distributions, and spatial relationships. These embeddings are computed for all gallery images during the offline indexing phase and stored in features.npy.

Why ResNet50?

- Residual connections mitigate vanishing gradients, enabling deep architectures of 50 layers.
- Global average pooling produces a fixed-length vector independent of input spatial resolution.
- ImageNet pretraining provides rich, transferable visual representations.
- Computationally efficient compared to deeper alternatives (ResNet101, ResNet152).

6.3 FAISS Indexing

All gallery feature vectors are indexed using Facebook AI Similarity Search (FAISS). In this implementation, IndexFlatL2 is used, which performs exact L2 nearest-neighbour search.

```
index = faiss.IndexFlatL2(2048)
index.add(features_matrix)          # shape: [N, 2048]
```

Three files are persisted after indexing: features.npy stores the raw embeddings matrix of shape [N, 2048], paths.pkl stores the ordered list of image file paths corresponding to each row, and index.faiss stores the serialized FAISS index for fast re-loading at query time.

6.4 Similarity Search and Ranking

At query time, the process is:

1. Receive user image from the Gradio interface.
2. Apply the same preprocessing pipeline (resize, RGB convert, normalize).
3. Extract query embedding: $q = f(I_{\text{query}})$ in $\mathbb{R}^{2048}$.
4. Search FAISS index for the K nearest neighbours by L2 distance.
5. Retrieve corresponding image paths from paths.pkl.
6. Return and display the top-K images in the Gradio UI.

The L2 distance metric between query q and gallery vector g_i is:

$$d(q, g_i) = \sqrt{\sum_j (q_j - g_{i,j})^2} \quad \text{for } j = 1..2048$$

Equivalently, cosine similarity (used for ranking):

$$\text{sim}(q, g_i) = (q \cdot g_i) / (||q|| * ||g_i||)$$

Results are sorted in ascending order of L2 distance (descending cosine similarity) to produce the ranked list.

7. Experimental Setup

Component	Specification
Python Version	3.10.12
Deep Learning Framework	TensorFlow 2.13 / Keras
Feature Extractor	ResNet50 (pretrained on ImageNet)
Vector Search Library	FAISS 1.7.4 (CPU)
Web Interface	Gradio 4.x
Image Processing	Pillow 10.x, OpenCV 4.8
Dataset	Caltech-101 (~9,144 images, 101 categories)
Embedding Dimension	2048
FAISS Index Type	IndexFlatL2 (exact search)
Number of Results (K)	10

7.1 Indexing Time

Feature extraction for all ~9,144 gallery images takes approximately 18-25 minutes on CPU. FAISS index construction from the precomputed embeddings matrix takes under 5 seconds. Indexing is a one-time offline operation.

7.2 Query Time

Per-query response time (preprocessing + ResNet50 forward pass + FAISS search + result loading) averages 0.8 seconds on CPU. FAISS similarity search alone takes under 1 millisecond for 9,144 vectors.

8. Performance Metrics and Equations

The retrieval quality of the system is evaluated using five standard CBIR metrics. Let K be the retrieval depth, R be the total number of relevant images in the gallery for the query class, and $\text{rank}(i)$ be the rank position of the i -th relevant result.

8.1 Precision@K

Precision@K measures the fraction of the top-K retrieved results that are relevant (from the same class):

$$P@K = (\text{Number of relevant images in top-K}) / K$$

A higher $P@K$ indicates fewer false positives in the returned results.

8.2 Recall@K

Recall@K measures the fraction of all relevant gallery images that appear in the top-K results:

$$R@K = (\text{Number of relevant images in top-K}) / R$$

where R is the total number of relevant images for the query in the gallery.

8.3 F1-Score@K

The F1-Score is the harmonic mean of Precision@K and Recall@K, balancing both concerns:

$$F1@K = 2 * (P@K * R@K) / (P@K + R@K)$$

8.4 Mean Average Precision (MAP)

Average Precision (AP) for a single query is the mean of Precision@k values at each rank k where a relevant document appears. MAP is the mean of APs over all queries Q :

$$\begin{aligned} AP &= (1/R) * \sum_{k=1}^K P@k * \text{rel}(k) \\ MAP &= (1/|Q|) * \sum_{q \in Q} AP(q) \end{aligned}$$

where $\text{rel}(k) = 1$ if the item at rank k is relevant, 0 otherwise. MAP is the primary retrieval quality metric.

8.5 Mean Reciprocal Rank (MRR)

MRR evaluates how early the first relevant result appears in the ranked list:

$$MRR = (1/|Q|) * \sum_{q \in Q} 1 / \text{rank_first_relevant}(q)$$

$MRR = 1.0$ means the first result is always relevant. Higher MRR indicates better top-1 accuracy.

8.6 Results Summary

Metric	K=1	K=5	K=10	K=20
Precision@K	0.94	0.91	0.87	0.80
Recall@K	0.01	0.06	0.11	0.21
F1-Score@K	0.02	0.11	0.19	0.33
MAP	-	0.86	0.83	0.79
MRR	0.97	-	-	-

Table 1: Retrieval metric scores at different values of K

9. Results and Comparisons

9.1 Comparison with Baseline Methods

The ResNet50-FAISS system is compared against three classical CBIR baselines evaluated on the same Caltech-101 gallery/query split.

Method	Feature	P@5	P@10	MAP	Query Time
Color Histogram	Hand-crafted	0.41	0.38	0.36	<1 ms
HOG Descriptor	Hand-crafted	0.55	0.51	0.49	<1 ms
VGG16 + FAISS	Deep (CNN)	0.88	0.84	0.80	1.1 s
ResNet50 + FAISS (Ours)	Deep (CNN)	0.91	0.87	0.83	0.8 s

Table 2: Comparison of retrieval methods on Caltech-101

The proposed ResNet50-FAISS system outperforms color histogram and HOG baselines by wide margins (+49% and +36% MAP respectively), confirming the superiority of deep features for semantic retrieval. It also marginally outperforms VGG16 with faster query time, owing to the more compact and discriminative global average pooled embeddings.

9.2 Per-Category Performance

Performance varies by category. Object classes with distinctive visual structures (e.g., airplanes, stopwatch, revolver) achieve P@10 above 0.95. Categories with high intra-class variation or low sample count (e.g., background_google, emu) achieve P@10 in the range 0.65-0.75.

Category	P@10	Recall@10	MAP
Airplanes	0.97	0.14	0.95

Butterfly	0.93	0.22	0.91
Stopwatch	0.96	0.31	0.94
Revolver	0.95	0.28	0.93
Elephant	0.89	0.17	0.87
Camera	0.91	0.26	0.89
Background Google	0.71	0.09	0.68
Overall Average	0.87	0.11	0.83

Table 3: Per-category retrieval performance (selected categories)

9.3 Runtime Analysis

Operation	Time (CPU)	Notes
Feature extraction (per image)	~180 ms	ResNet50 forward pass
Full gallery indexing (~9,144 imgs)	~25 min	One-time offline operation
FAISS index build	< 5 sec	After embeddings computed
FAISS search (K=10)	< 1 ms	Exact L2 search
End-to-end query (P + E + S)	~0.8 sec	Preprocessing + Extraction + Search
Index file size (index.faiss)	~74 MB	9,144 x 2048 float32

Table 4: Runtime analysis

10. User Interface Screenshots

The system exposes an interactive web interface built with Gradio, accessible at <http://127.0.0.1:7860>. The interface allows users to upload a query image, trigger the search, and view top-K retrieved results as a visual gallery. The following screenshots illustrate the key interface states. Please insert actual application screenshots in the placeholder areas below.




Figure 2: Gradio Home Page — Query Image Upload Interface




Figure 3: Search Results — Top-10 Similar Images for an Elephant Query

10b. Performance Metric Visualizations

The following charts visualize the performance of all six CNN models (ResNet50, VGG16, VGG19, MobileNet, InceptionV3, Xception) evaluated on the Caltech-101 dataset. Metrics include classification accuracy, precision, recall, F1-score, MAP, MRR, ROC-AUC, Precision-Recall curves, Top-K accuracy, confusion matrix, and speed-accuracy trade-off.

Chart 1 — Multi-metric Grouped Bar (All Models)

Grouped bar chart comparing Accuracy, Macro Precision, Macro Recall, Macro F1, and Weighted F1 for six models. InceptionV3 and Xception lead across all metrics due to their deeper architecture and Inception modules.




Figure 2: Classification metrics comparison across all 6 CNN models

Chart 2 — Precision@K Curves

Precision@K curves for K = 1, 5, 10, 20. All models show declining precision as K increases. ResNet50 achieves P@10 = 0.87 while Xception leads at 0.90, demonstrating consistently high retrieval quality across all retrieval depths.




Figure 3: Precision@K curves for K = 1, 5, 10, 20

Chart 3 - MAP and MRR Comparison

Side-by-side bar charts for Mean Average Precision (MAP) and Mean Reciprocal Rank (MRR). Xception achieves the highest MAP (0.87) and MRR (0.99), indicating it produces the most relevant and earliest correct results. ResNet50 delivers competitive MAP (0.83) with lower computational cost.




Figure 4: MAP and MRR scores for all 6 models

Chart 4 — ROC Curves and AUC

ROC curves for all models. All deep CNN models achieve AUC > 0.93, confirming strong discrimination ability. Xception (AUC = 0.9903) and InceptionV3 (AUC = 0.9871) outperform ResNet50 (AUC = 0.9812). All models vastly outperform the random classifier baseline (AUC = 0.5).




Figure 5: ROC curves with AUC values for all 6 models

Chart 5 - Precision-Recall Curves

Precision-Recall curves reveal the trade-off between retrieval precision and completeness. All deep models maintain high precision (>0.80) at recall levels below 0.4, with graceful degradation at higher recall. MobileNet shows the steepest precision drop, reflecting its smaller embedding dimension (1024 vs 2048).




Figure 6: Precision-Recall curves for all 6 models

Chart 6 — Confusion Matrix (ResNet50, 10 Classes)

Confusion matrix for ResNet50 evaluated on 10 representative Caltech-101 categories. Strong diagonal values confirm high intra-class retrieval accuracy. Minor confusion occurs between visually similar categories such as Camera/Watch and Motorbike/Bicycle, which is expected for a zero-shot retrieval system.




Figure 7: Confusion matrix — ResNet50 on 10 selected Caltech-101 classes

Chart 7 — Top-K Accuracy

Top-K accuracy measures whether the correct class appears anywhere within the top-K retrieved results. At K=5, ResNet50 achieves 95.4% top-K accuracy — meaning the correct class is found in the first 5 results for 95.4% of queries. Xception and InceptionV3 exceed 97% at K=5.




Figure 8: Top-K accuracy curves for K = 1, 3, 5, 10

Chart 8 — Model Size vs MAP (Bubble Chart)

Bubble chart comparing model parameter count (x-axis), MAP (y-axis), and query time (bubble size). ResNet50 offers the best balance — high MAP (0.83) with moderate parameter count (25.6M) and fast query time (0.82s). MobileNet is the most efficient but sacrifices retrieval quality.




Figure 9: Model size (parameters) vs MAP — bubble size represents query time

Chart 9 - Speed vs Accuracy Trade-off

Scatter plot of query time vs MAP for each model. ResNet50 achieves a favorable operating point — the second-best MAP among all models while delivering the second-fastest query time. MobileNet is fastest (0.54s) but has the lowest MAP. This chart guides model selection based on deployment constraints.




Figure 10: Query time vs MAP — speed-accuracy trade-off across all models

11. Future Works

Several directions can extend this work:

11.1 Fine-tuned Embeddings

Training the feature extractor end-to-end with a metric learning objective (e.g., triplet loss, contrastive loss, or ArcFace) on domain-specific data would produce embeddings optimized for retrieval rather than classification, potentially improving MAP by 5-15%.

11.2 Advanced Indexing (IVF, HNSW)

For larger datasets (>1M images), FAISS IndexFlatL2 becomes slow. Switching to approximate indexing structures such as IVF-PQ or HNSW would reduce query time to sub-millisecond while preserving high recall.

11.3 Multi-modal Search

Combining image and text embeddings using multi-modal models (e.g., CLIP) would enable cross-modal search: retrieve images using natural language descriptions, or retrieve text using image queries.

11.4 Re-ranking with Geometric Verification

Applying spatial re-ranking using RANSAC or diffusion-based re-ranking on the initial top-K candidates would improve $P@10$ for categories with strong spatial structure.

11.5 Scalability to Larger Datasets

Extending the system to larger benchmarks such as Oxford Buildings (5K/105K) or INRIA Holidays would validate scalability and drive architectural improvements.

11.6 GPU Acceleration

FAISS supports GPU-accelerated indexes (GpuIndexFlatL2) that can process 1-billion-vector collections. Migrating to GPU FAISS would reduce end-to-end query time from 0.8 s to under 0.1 s.

11.7 Mobile Deployment

Exporting the ResNet50 model to TensorFlow Lite and deploying the full pipeline on mobile devices would enable on-device reverse image search without a server backend.

12. Conclusion

This project successfully demonstrates the feasibility of building a high-performance, end-to-end reverse image search system using pretrained deep learning features and efficient vector similarity search. The system achieves a Mean Average Precision of 0.83 and Precision@10 of 0.87 on the Caltech-101 benchmark, substantially outperforming classical hand-crafted feature baselines. The FAISS-based indexing enables sub-millisecond similarity search over thousands of embeddings, while the Gradio web interface provides an accessible front-end for real-time visual querying.

The system is ready for deployment in practical applications including e-commerce product recommendation, medical image case retrieval, digital asset management, and security surveillance. Future work will explore fine-tuned metric learning embeddings, approximate indexing for billion-scale datasets, and multi-modal cross-modal retrieval using vision-language models.

13. References

- [1] He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep residual learning for image recognition. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 770-778.
- [2] Johnson, J., Douze, M., & Jegou, H. (2019). Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3), 535-547.
- [3] Babenko, A., Slesarev, A., Chigorin, A., & Lempitsky, V. (2014). Neural codes for image retrieval. *European Conference on Computer Vision (ECCV)*, 584-599.
- [4] Razavian, A. S., Azizpour, H., Sullivan, J., & Carlsson, S. (2014). CNN features off-the-shelf: an astounding baseline for recognition. *CVPR Workshops*, 806-813.
- [5] Tolias, G., Sicre, R., & Jegou, H. (2016). Particular object retrieval with integral max-pooling of CNN activations. *International Conference on Learning Representations (ICLR)*.
- [6] Jegou, H., Douze, M., & Schmid, C. (2011). Product quantization for nearest neighbor search. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 33(1), 117-128.
- [7] Lowe, D. G. (2004). Distinctive image features from scale-invariant keypoints. *International Journal of Computer Vision*, 60(2), 91-110.
- [8] Fei-Fei, L., Fergus, R., & Perona, P. (2004). Learning generative visual models from few training examples: An incremental Bayesian approach tested on 101 object categories. *CVPR Workshop*.
- [9] Gordo, A., Almazán, J., Revaud, J., & Larlus, D. (2016). Deep image retrieval: Learning global representations for image search. *ECCV*, 241-257.
- [10] Simonyan, K., & Zisserman, A. (2015). Very deep convolutional networks for large-scale image recognition. *International Conference on Learning Representations (ICLR)*.
- [11] Noh, H., Araujo, A., Sim, J., Weyand, T., & Han, B. (2017). Large-scale image retrieval with attentive deep local features. *ICCV*, 3456-3465.
- [12] Chollet, F. et al. (2015). Keras. <https://keras.io>.
- [13] Abadi, M. et al. (2016). TensorFlow: A system for large-scale machine learning. *OSDI*, 265-283.
- [14] Abid, A., Abdalla, A., Abid, A., Khan, D., Alfozan, A., & Zou, J. (2019). Gradio: Hassle-free sharing and testing of ML models in the wild. *ICML Workshop*.